


FOUNDING PRODUCT DESIGNER • 10-MONTH CASE STUDY

# Computis

Building professional trust into an AI tax engine. From a v1 that 6 of 6 CPAs rejected, to a graduated-autonomy product that became the wedge.

Includes 2 misses I'd answer for. Each one generated a principle the deck still uses.

- P1 • TRUST-FIRST
- P2 • TRANSPARENCY
- P3 • AUGMENTATION



• TRANSACTIONS SURFACE • DAILY-DRIVER • CONFIDENCE AS FIRST-CLASS UI • ANOMALY REVIEW AS BOUNDED SURFACE

SECTION 01 — CONTEXT

# What I walked into.

A platform with technical depth and an AI feature set chasing the wrong problem.  
Compliance posture under two regimes. Eight weeks to pilot.

01

## Context — company, product, users, compliance.

### The company

Seed-stage B2B SaaS. Pre-Sean: technical demo with strong AI, no daily-driver design. Team of 6.

### The product

Crypto-tax preparation platform for professional accounting firms. Ingest wallet/exchange data → categorize → produce IRS-ready forms.

### The users

Senior CPAs — license on the line; skeptical of automation. Tax analysts — high-volume, override-heavy operators inside firms.

### Compliance posture — two regimes, both load-bearing

#### REGIME A

##### IRS audit defensibility

Every classification must trace to a rule, a transaction hash, and a timestamp. "Defend it to the IRS" is the user-quoted bar.

#### REGIME B

##### WCAG 2.1 AA

Tax-prep is a regulated profession with accessibility obligations. Color-only encoding is disallowed; all tier states ship with text + shape.

*Both regimes shaped the design system — see Slide 12.*

### My team

CEO/CTO (sets product direction, reviews architecture). 2 engineers (one full-stack, one ML-leaning). Me. **No PM, no other designer.** Decisions documented; principles doc as tiebreaker.

# The CPA's day. Where Computis intervenes.

95% of time on manual classification today. Computis targets the inverse — 85% strategic advisory, 15% review.

BEFORE COMPUTIS · TODAY

**95%** manual classification & error hunting

**5%** strategic tax advisory

6+ hrs/day manual classification    High audit-liability exposure

AFTER COMPUTIS · TARGET STATE

**85%** strategic tax advisory & client service

**15%** review queue — flagged items only

Focus on high-value advisory    IRS-ready audit-trail protection

The inversion is the design brief. Every surface in this deck is judged against whether it moves the 95/5 toward 15/85 — without the CPA losing the audit-defensibility their license depends on.

Baseline: CPA shadow sessions, weeks 4-6 · target state validated against beta cohort

# The problem, sized.

## User problem

Crypto-tax prep takes **4–8 hours per client** at firms — most of it spent verifying AI-classified transactions one-by-one because **none of the classifications come with a reason**.

*A CPA who can't explain a number won't sign it.*

## Business problem

Trial → paid conversion stalled. Demo CPAs liked the speed; pilot CPAs rejected the product after one client load.

Sources: 12 semi-structured CPA interviews · 6 shadow sessions · 50K transactions sampled across 4 pilot firms · 300+ G2 reviews of competing products coded for themes.

### The numbers I was working from

4–8h

PER-CLIENT PREP TIME, BASELINE

2h

OF THAT, ANOMALY INVESTIGATION

~1.5h

OF THAT, CLASSIFICATION REVIEW

73%

TRANSACTIONS FOLLOWING PREDICTABLE PATTERNS

73% predictability **is** the wedge: automate the predictable transparently, give CPAs leverage on the rest.

*"The wedge is in **how the AI behaves**, not **what it does**."*

WORKING HYPOTHESIS · WEEK 3 · VALIDATED WEEK 11

SECTION 02 — RESEARCH

# Three findings broke our v1 in week 7.

12 interviews · 6 shadow sessions · 50K transactions · 300+ G2 reviews

02

# What 12 CPAs told us. Four quotes carry the argument.

Semi-structured interviews, weeks 1–6, four firms. Coded for themes. These four quotes carried 80% of the thematic weight.

QUOTE • 01 • TIME ALLOCATION

*"I spend 6 hours classifying before I can do any actual tax strategy work."*

CPA, 15+ years

6+ HRS/DAY

QUOTE • 02 • AUDIT DEFENSIBILITY

*"If the IRS asks why I classified something as a staking reward, I need to show my reasoning."*

Tax Partner, Big 4

CITED 11/12

QUOTE • 03 • LATE ERROR DISCOVERY

*"I found an FMV error after we'd already sent the return. We had to file an amendment."*

Senior Tax Manager

4 OF 12 HIT

QUOTE • 04 • ENGINEERING DEPENDENCY

*"Every rule change requires an engineering ticket. I wait 3–5 days for simple pattern updates."*

Crypto Tax Specialist

3–5 DAY WAIT

**Synthesis →** Automate the 73% **transparently**. Give CPAs leverage on the 27%. Treat overrides as additional context, not corrections. **This became the principles doc.**

Source: semi-structured interviews • weeks 1–6 • 4 firms • n=12 • coded for themes

# How we worked — research in four phases, project across ten months.

Cadence held: two-week design, one-week eng-feedback. The v1→v2 pivot fit inside the existing rhythm.



## 0 of 6 rejected our V1.

Full automation. The model worked. The product did not.

### V1 • Full automation prototype • Week 7 pilot

V1 • REJECTED

Auto-classify all transactions. Surface result. Allow override if the CPA disagrees. Optimized for speed; the model was acceptable.

*"The AI part is fine. The black box part is not."*

SENIOR CPA • WEEK 7 PILOT INTERVIEW • CITED IN 11 OF 12 INTERVIEWS

0 / 6

CPAs in the V1 pilot would load a second client. Six of six.

PILOT TEST • WEEK 7 • BLIND PROTOCOL

### What we initially thought was wrong

- Model accuracy — fix the classifier.
- Speed of override — let CPAs correct faster.
- Result presentation — better confidence display.

#### WHAT WE WERE WRONG ABOUT

All three. None of those was the rejection. The model was acceptable. **Speed was never the objection. Opacity was.**

→ Next slide: what the rejection actually said.

# What the rejection actually said.

Overrides weren't corrections. They were context the AI couldn't see. The reframe became V2.

## V2 · Graduated autonomy · Week 11 pilot

V2 · ADOPTED

AI tiers itself by confidence. CPA controls the routing. Override flow treats new input as *context*, not correction.

*Same 6 CPAs · same blind protocol · 4 weeks later.*

PILOT TEST · WEEK 11 · FIRMS-MATCHED COHORT

6 / 6 CPAs loaded a second client in the V2 pilot.

23% → 89%

CLASSIFICATION TRUST RATE  
LIKERT · PRE/POST V2 · SOURCE: PILOT SURVEY · n=6 · WK 7 vs WK 11

↓85%

MANUAL CLASSIFICATION RATE  
ENG TELEMETRY · ≠ ENG-DEP ↓85%

## What changed in 4 weeks

- Confidence tiers became a first-class UI element (slide 13).
- Every classification ships with a *why* — source rule, similar decisions, confidence breakdown.
- Override flow treats new input as context, not correction.

## What didn't change

- The underlying classifier. Same model. Different surfaces.
- The 73% predictability number — same data sample, same finding.

The pivot wasn't AI quality. It was AI behavior. Same predictions. Different presentation. Opposite outcome.

# How I got the team there — the joint-session move.

CTO · WEEK 7

"The model accuracy is 89%. That's the metric. The override behavior is a UI tuning problem."

SEAN · WEEK 7

"The CPAs aren't overriding because the model is wrong. They're overriding because the model can't explain itself."

CTO · WEEK 8 · AFTER JOINT SESSION

"Overrides aren't corrections. They're additional context the AI doesn't have."

## Joint session · Week 8 · Synthesis board

CPA OVERRIDE

"jurisdictional context"

AI OUTPUT

"pattern + rule"

JOINT FRAME

override = context AI couldn't see

"I'm not correcting the AI. I'm adding what it can't see — that this client is a partnership, not an LLC."

— David Chen · session 4 · pilot firm B

DECISION · TREAT OVERRIDES AS CONTEXT

P3 Augmentation · principle locked · week 8

RECONSTRUCTION OF FIGJAM SYNTHESIS BOARD · SESSION WITH CEO/CTO + 2 SENIOR CPAS · 2.5 HOURS

SECTION 03 — DESIGN SYSTEM & UI CRAFT

# Where the system became a product.

Three principles. One token foundation. Three surfaces that earned the metrics.

03

## Three principles. Every decision traces back.

### P1 • TRUST-FIRST

#### Admit uncertainty.

- **Decision:** Every classification ships with a confidence tier. *Tradeoff:* +cognitive load on the row.
- **Decision:** Below-60% tier surfaces as "Review needed", not a guess. *Tradeoff:* manual share visible (~15%).
- **Decision:** Confidence is a first-class UI element, not a hover-state.

### P2 • TRANSPARENCY

#### Every AI decision has a "why."

- **Decision:** Drawer (not modal) preserves row context. *Tradeoff:* narrower content column.
- **Decision:** Source rule + similar decisions + confidence breakdown rendered, not linked.
- **Decision:** "Defend it to the IRS" is the bar; every artifact must trace.

### P3 • AUGMENTATION


#### AI suggests. CPA decides.

- **Decision:** Overrides recorded as context, not correction; AI learns within a firm.
- **Decision:** No auto-action on Low-tier without explicit CPA acknowledgement.
- **Decision:** Rule Builder is CPA-authored, not engineer-authored.

EVERY DECISION IN THIS DECK ROUTES THROUGH ONE OF THESE THREE. WHEN TWO PULL AGAINST EACH OTHER, THE PRINCIPLES DOC IS THE TIEBREAKER.

# The solution — 3-tier AI confidence system.

Confidence is a first-class UI element. The tier determines the workflow.




# Why three tiers, and why these cutoffs.

Four alternatives tested. Three rejected. The cheapest way to make a confidence number behave like a routing instruction.

THRESHOLDS CONSIDERED

|                                                                                                                                               |          |
|-----------------------------------------------------------------------------------------------------------------------------------------------|----------|
| <b>5-tier</b><br>CPAs flattened it mentally — "top two are auto, bottom three are manual." Granularity didn't map to workflow.                | REJECTED |
| <b>2-tier</b><br>Collapsing "review" into "manual" lost the routing affordance. The Medium-tier behavior is neither auto nor manual.          | REJECTED |
| <b>Percentile-based</b><br>Thresholds drifted as the model improved. CPAs lost calibration — "Medium" meant something different week to week. | REJECTED |
| <b>3-tier fixed-band</b><br>≥90 auto · 60–89 review · <60 manual. Fixed bands give CPAs a stable mental model that survives model drift.      | ADOPTED  |

VALIDATION — HIGH-TIER ACCEPTANCE ACROSS ITERATIONS




Iteration 3 — chosen cutoffs (≥90 / 60–89 / <60) tested against same 50K-tx sample. Acceptance arc shape reproducible; specific bar heights reconstructed from pilot survey data and engineering telemetry.

Thresholds aren't arbitrary. They're the cheapest way to make a confidence number behave like a routing instruction.

# Visual Rule Builder — no-code rule authoring.

Single biggest contributor to **↓85% engineering dependency**. Pre-Builder: every rule = code change. Post-Builder: CPA-authored, version-controlled.

• RULE ENGINE • ACTIVE RULES + CONFLICT DETECTION • AUDIT LOGGED



Production capture • 5 active rules • classification + priority conflict detection • per-rule affected-txn count

Active Rules surface — the management view CPAs land in. Next slide: the rule-authoring surface (Create Merge Rule modal).

## Pre-save preview

Every rule shows **what it would do** against the last 30 days before activate. CPAs see the impact before committing.


## Why this earned the 85%

Pre-Builder, every rule was a PR. Engineering wrote, deployed, tested — 12-week PR cadence. Post-Builder, CPAs author. Engineering builds the primitives. **85% fewer rule-change PRs**, 12 weeks pre vs 12 weeks post.

# Rule authoring — conditions, actions, preview-before-save.

The companion surface to the management view. **Preview-before-save** is the fix to the Slide 18 miss-card — CPAs cannot save a rule without seeing its impact.

● CREATE MERGE RULE · CONDITIONS · ACTIONS · PREVIEW IMPACT



Production capture · Create Merge Rule modal · BTC Mining Income rule in preview state · 23-tx impact estimate

Paired with previous slide's management view. Together they cover the full Rule Builder lifecycle: manage existing rules → author a new one → preview before save.

## Conditions

Three AND-joined predicates building the rule's match logic. **chip-filter** for values, dropdown for operators, AND/OR toggle on the logical operator. Same primitives as the Transactions filter bar.

## Actions

Three chained classifications — Transfer / Mining + Income / Pending — applied in order when the conditions match. **tag-category chips** reused from the classification badges on Transactions.


## Preview Impact

**Non-optional preview.** 23 transactions will be affected · Before/After classification breakdown shown · cannot save without acknowledging impact. This is the Slide 18 miss-card, fixed.

**Component-driven.** Drawer · chip-filter · tag-category · before/after card — all reused. Why this shipped in 6 weeks, not 12.

# Information architecture — where every AI feature lives.

AI is concentrated in **Transactions** and a few tagged settings — never on Clients or Reports, where the CPA's authority must read as unambiguous.



**Reading the IA:** 9 AI-tagged features total. 4 sit inside Transactions — by design. The CPA's authority surfaces (Clients, Reports) carry exactly 3 AI features between them, each scoped to non-classification work.

Source: Figma file computis-ia-v2 · 5 top-level domains · 9 AI-tagged features · governed by P3 (augmentation) principle

# Token foundation — type, color, spacing, components.

A system gets enforced or it doesn't exist. Tokens named by intent, not by size. 47 components in production resolve to these primitives.

01 · TYPE SCALE

Display68 / 1.05

H1 — Slide title48 / 1.15

H2 — Card title26 / 1.25

Body — 20 / 1.50Noto Sans

EYEBROW · MONOJetBrains

02 · COLOR — NAMED BY INTENT

BRAND

--navy  
#0A2540

--p1-teal  
#00D4AA

--p2-blue  
#2E86C1

--p3-amber  
#F5A623

CONFIDENCE TIERS — TEXT + COLOR, NEVER COLOR ALONE

HIGH--tier-high

MED--tier-med

LOW--tier-low

03 · SPACING — 4PX BASE

--space-14px

--space-28px

--space-416px

--space-624px

--space-732px

--space-848px

--space-964px

04 · COMPONENT LIBRARY — SAMPLED

BUTTON

Primary action

Secondary

Ghost link →

CONFIDENCE BADGE

HIGH · 94%

MED · 76%

LOW · 41%

INPUT & CHIP-FILTER

Search transactions...

0x91dd...44ab

AllHighMed

DATA CARD

CLASSIFICATION TRUST

89%

+11pts vs v1 pilot

Source: tokens.css + Figma component library · authored wk 10 · 47 production components resolve to these primitives

SECTION 03 — DESIGN SYSTEM · TOKENS

# WCAG 2.1 AA + IRS audit defensibility — regulatory constraints made the IA tighter, not looser.

REGIME A · WCAG 2.1 AA

## Four decisions, made at the token level.

DECISION · 01

**Color + label, never color alone.** Confidence tier chips pair color encoding with text labels (High / Medium / Low). Color-blind safe; AT-readable.

DECISION · 02

**focus-visible on every interactive surface.** Audit drawer has 8 tab stops, all visible. No mouse-only affordances.

DECISION · 03

**4.5:1 minimum contrast at the token level.** Verified in `tokens.css` , not slide-by-slide. SOURCE: `tokens.css` audit · wk 18

DECISION · 04

**Reduced-motion variant for drawer slide-in.** Vestibular triggers in a tax-prep context aren't hypothetical. SOURCE: `tokens.css` audit · wk 18

REGIME B · IRS AUDIT DEFENSIBILITY

## Four decisions, embedded in product surface.

DECISION · 01

**Timestamped audit trail per classification.** Embedded in transaction context, not a separate log.

DECISION · 02

**Override input requires reason taxonomy.** Not free-text. Free-text doesn't hold up at audit.

DECISION · 03

**Provenance metadata in every export.** Who · when · what tier · what override. Subpoena-ready by default.

DECISION · 04


**The "defend it to the IRS" bar.** Every drawer surface reads as if it could be subpoenaed verbatim. Because some will.

The two regimes didn't compete for canvas. WCAG forced label semantics. IRS forced data semantics. The IA absorbed both without surface bloat — and the IA diagram (already at slide 11) is what made the absorption possible.

# Sketch → Mid-Fi → Production. The Transactions surface across 16 weeks.


UI DESIGN EVOLUTION · SKETCH · MID-FIDELITY · PRODUCTION

## Transactions Screen — Iteration Progression




### 01 — Sketch

Information architecture & screen layout exploration



### 02 — Mid-Fidelity

Component structure, data hierarchy & states



### 03 — Production UI

Token system, AI confidence patterns, dark sidebar

#### 01 Reviewable suggestions

Status states surface uncertainty before filings

#### 02 Confidence indicators

Explainable AI pattern—never a black box

#### 03 AI Classification chips

Color-coded transaction types for instant scan

#### 04 Anomaly Flags surface

Critical issues elevated to the primary review

DISCOVERY · WK 1

BUILD · WK 6

SHIP · WK 16

OUTCOMES · 45% reduction in client onboarding · 32% increase in demo-to-conversion · 85% reduction in engineering dependency

Authored weeks 1-16 · the visible record of iteration on the Transactions surface · outcomes annotated in production state

## Two columns, one strip — the AI delta on the Transactions surface.

• BEFORE • WITHOUT AI • 6 COLUMNS

| <input type="checkbox"/> | View | Date ↑↓    | Type    | Asset           | Amount   | FMV (USD)  | Status    | Actions |
|--------------------------|------|------------|---------|-----------------|----------|------------|-----------|---------|
| <input type="checkbox"/> |      | 2022-06-15 | Receive | Bitcoin (BTC)   | 0.25 BTC | \$5,250.00 | Confirmed | ✓       |
| <input type="checkbox"/> |      | 2022-07-02 | Swap    | Ethereum (ETH)  | 1.5 ETH  | \$2,850.00 | Suggested | ✓       |
| <input type="checkbox"/> |      | 2022-08-14 | Send    | Bitcoin (BTC)   | 0.15 BTC | \$3,750.00 | Flagged   | ✓       |
| <input type="checkbox"/> |      | 2022-09-05 | Receive | USD Coin (USDC) | 500 USDC | \$500.00   | Pending   | ✓       |
| <input type="checkbox"/> |      | 2022-10-22 | Merge   | Ethereum (ETH)  | 0.75 ETH | \$1,125.00 | Confirmed | ✓       |
| <input type="checkbox"/> |      | 2022-11-08 | Swap    | Bitcoin (BTC)   | 0.08 BTC | \$1,050.00 | Suggested | ✓       |
| <input type="checkbox"/> |      | 2022-12-01 | Receive | Solana (SOL)    | 10 SOL   | \$350.00   | Flagged   | ✓       |

Showing 1 of 7 transactions

Bulk Actions: Accept All Tag Export

### 1 AI CLASSIFICATION COLUMN

New column: AI's best guess at transaction type — Income, Trade, Expense, Transfer. P2 (transparency): the read-only label tells the CPA what the AI thinks before they look at confidence.

### 2 CONFIDENCE COLUMN

New column: tier-aware percentage, click-through to audit drawer. P2 again, on a different axis: *how sure* separate from *what*. Drives the 3-tier workflow visible on Slide 13.

• AFTER • WITH AI • +2 COLUMNS + ANOMALY STRIP

| Confidence: <span>AI</span> <span>High</span> <span>Medium</span> <span>Low</span> Status: <span>AI</span> <span>Confirmed</span> <span>Suggested</span> <span>Flagged</span> <span>Filters</span> <span>Bulk Tag</span> <span>AI Classify</span> |      |            |         |                 |          |            |                   |            |           |         |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------|------------|---------|-----------------|----------|------------|-------------------|------------|-----------|---------|
| <input type="checkbox"/>                                                                                                                                                                                                                          | View | Date ↑↓    | Type    | Asset           | Amount   | FMV (USD)  | AI Classification | Confidence | Status    | Actions |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-06-15 | Receive | Bitcoin (BTC)   | 0.25 BTC | \$5,250.00 | Income            | 65%        | Confirmed | ✓       |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-07-02 | Swap    | Ethereum (ETH)  | 1.5 ETH  | \$2,850.00 | Trade             | 78%        | Suggested | ✓       |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-08-14 | Send    | Bitcoin (BTC)   | 0.15 BTC | \$3,750.00 | Expense           | 45%        | Flagged   | ✓       |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-09-05 | Receive | USD Coin (USDC) | 500 USDC | \$500.00   | Unclassified      | 0%         | Pending   | ✓       |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-10-22 | Merge   | Ethereum (ETH)  | 0.75 ETH | \$1,125.00 | Transfer          | 92%        | Confirmed | ✓       |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-11-08 | Swap    | Bitcoin (BTC)   | 0.08 BTC | \$1,050.00 | Trade             | 82%        | Suggested | ✓       |
| <input type="checkbox"/>                                                                                                                                                                                                                          |      | 2022-12-01 | Receive | Solana (SOL)    | 10 SOL   | \$350.00   | Income            | 55%        | Flagged   | ✓       |

Showing 1 of 7 transactions

Bulk Actions: Accept All Tag Export

#### Anomaly Flags

Issues requiring attention

Volume Spike

Unusual Activity

Count: 15

Status: High

Investigate →

Missing FMV

Pricing Requested

Count: 8

Status: Critical

Fix Values →

Rule Conflict

Classification Issue

Count: 2

Status: Low Priority

Resolve →

#### AI Classification

Confidence distribution

High

61 items

55%

Medium

40 items

34%

Low

14 items

11%

Total Classified: 117 transactions

### 3 ANOMALY FLAGS STRIP

Below the table. Volume Spike · Missing FMV · Rule Conflict surfaced proactively. P3 (augmentation): AI choosing what the CPA should look at next — not just classifying when asked.

SECTION 03 — UI CRAFT • BEFORE / AFTER

# Transactions — the daily-driver surface.

TRANSACTIONS · 1,247 RESULTS · Q2 2024

[1] DENSITY · 14 ROWS

| DATE  | TX            | FROM → TO           | USD         | CATEGORY <sup>2</sup> | CONF <sup>3</sup> |
|-------|---------------|---------------------|-------------|-----------------------|-------------------|
| 04/01 | 0xa3f2...d891 | Coinbase → Wallet 2 | \$4,210.55  | Transfer              | 94%               |
| 04/02 | 0xbc14...7e02 | Wallet 2 → Kraken   | \$1,820.00  | Trade                 | 91%               |
| 04/03 | 0x91dd...44ab | Unknown → Wallet 1  | \$612.30    | Income                | 76%               |
| 04/05 | 0x71f8...0c2e | Wallet 1 → ?        | \$98.40     | Review needed         | 52%               |
| 04/06 | 0xab90...22f1 | Metamask → DEX      | \$2,003.10  | Swap                  | 93%               |
| 04/07 | 0x44b2...1aa1 | Wallet 3 → Coinbase | \$502.00    | Transfer              | 96%               |
| 04/08 | 0xcc09...8bf3 | Binance → Wallet 1  | \$11,500.00 | Trade                 | 92%               |
| 04/09 | 0x10de...44ce | Wallet 1 → Wallet 2 | \$240.00    | Transfer              | 81%               |
| 04/10 | 0xee21...f0a0 | Coinbase → Wallet 2 | \$3,005.05  | Income                | 94%               |
| 04/11 | 0x88c4...dd71 | Wallet 2 → ?        | \$48.00     | Review needed         | 41%               |
| 04/12 | 0x52a9...b302 | Wallet 2 → Kraken   | \$908.40    | Trade                 | 95%               |
| 04/13 | 0xfd00...3417 | Metamask → DEX      | \$721.10    | Swap                  | 91%               |
| 04/15 | 0xa120...2274 | Wallet 1 → Wallet 3 | \$60.00     | Transfer              | 74%               |
| 04/16 | 0x09bc...ee55 | Coinbase → Wallet 2 | \$1,200.00  | Income                | 90%               |

1-14 OF 1,247

BULK-SELECT <sup>5</sup> · DISABLED UNTIL CONFIRM-TIER-BY-TIER

- 1

Density · 14 rows visible. Tax analysts work in batches of 20–40. Below 14, the rhythm breaks.
- 2

Category column shows AI suggestion + "Review needed" for Low-tier. Never a guess on Low.
- 3

Confidence badge inline · WCAG 2.1 AA. Tier label + color + position; not color-only.
- 4

Row tint mirrors tier. Subtle (4% alpha) — visible at scan, doesn't dominate.
- 5

Bulk-select shipped last. See callout — principle-tradeoff signature.

WHY BULK-SELECT SHIPPED LAST

We shipped it last because defaulting to bulk-override would have unraveled the per-row trust signal we'd just built. P1 ↔ P3 tension; P1 wins on bulk by default. Bulk is opt-in, gated, and explicitly logged.

# Audit trail — every state change timestamped, attributed, exportable.

AUDIT TRAIL · EXPORT-LEVEL PROVENANCE · TIMESTAMPED · EXPORTABLE

Search...

Export

Generate IRS 8949, QBO, and CSV files with embedded audit logs

Tax Year: 2023 2022 2021 View Audit Trail

Recent Exports

Download and manage your export history

| Date         | Export Type | Tax Year | Transactions | Status   | User      |
|--------------|-------------|----------|--------------|----------|-----------|
| Oct 24, 2023 | CSV         | 2022     | 124          | Complete | John Smit |
| Oct 20, 2023 | IRS 8949    | 2022     | 120          | Complete | John Smit |
| Oct 15, 2023 | QBO         | 2022     | 118          | Complete | John Smit |

IRS 8949 Export

Precisely curated gains report

Transactions Ready: 24

Unclassified Transactions: 12

Missing Fmv Data: 4

AI Classified Transactions: 8

Export 8949

QuickBooks Export

QBO format for accounting

Transactions Ready: 124

Qbo Mapping Status: Complete

Chart Of Accounts: Standard Crypto

Last Exported: Never

Export QBO

CSV Export

Raw data with audit tra

Transactions Ready:

Export Format:

Include Audit Data:

Data Validation:

Export Configuration

Customize your export settings

Data Inclusion

Select which data to include in exports

☒ Include all transactions

☒ Include confirmed transactions only

☒ Include audit trail data

☒ Include AI confidence scores

☒ Include FMV source data

Data Validation

Your data has been validated for export

Transaction Classification

All classifications validated

FMV Data Accuracy

Some pricing data missing

AI Classification Confidence

Audit Trail

2023-10-24 14:23:45

Export Initiated

john.smith@company.com

CSV export started for Tax Year 2022 (124 transactions selected)

2023-10-24 14:23:46

Data Validation

System

All transactions validated: FMV data complete, cost basis verified

2023-10-24 14:24:12

AI Classification Review

System

AI confidence scores calculated: 118 High, 6 Medium confidence classifications

2023-10-24 14:24:38

Manual Review

john.smith@company.com

User reviewed 6 medium-confidence transactions and approved classifications

2023-10-24 14:25:01

CSV Generation

System

CSV file generated with all required fields and audit trail metadata

2023-10-24 14:26:15

Export Approved

john.smith@company.com

Final export approved and marked for download (File: crypto\_export\_2022.csv)

2023-10-24 16:45:22

Export Trail

Print

## Why audit trail, not just logs

Logs answer *what happened*. The audit trail answers **who did it, when, in what order, and with what input** — the four questions the IRS asks. Every export carries this trail. Tradeoff: more state to render. Worth it.

## Reason taxonomy — what CPAs actually chose


|                    |     |
|--------------------|-----|
| Wrong category     | 52% |
| Missing context    | 23% |
| Firm-specific rule | 18% |
| Jurisdictional     | 5%  |
| Other              | 2%  |

SOURCE: pilot override telemetry · n=~2,100 · 90d window

Production capture · 5 events on a single export · CPA + system attribution interleaved · exportable as PDF or signed log

# Smart wallet upload — the onboarding wedge.

WALLET INGESTION · UPLOAD → VALIDATE → MAP → IMPORT



Production capture · 4-step ingestion · format auto-detect · per-file validation state surfaced before mapping

## Problem

V1 onboarding asked the CPA to map columns from raw CSV → schema, blind. **38% drop-off** — at step 2.

## Redesign

Auto-detect source exchange from header signature. Propose a mapping with per-column confidence. CPA reviews, doesn't author.

## Outcome

↓**45% onboarding time**, v1 median 14.5d → v2 median 8.0d (Slide 17). Drop-off at step 2 fell into single digits.

### P3 · AUGMENTATION

Don't replace the CPA's judgment about what file came from where. Accelerate it. *Detect*, then *propose*, then *confirm*.

# Four verified metrics. Two specific misses I'd answer for.

Each metric sourced and caveated. The misses aren't roadmap — they're decisions I'd reverse.

04

# How design + 2 engineers shifted classification-logic ownership.

Pair-design in PRs · principles doc as tiebreaker · async handoff in Linear · sync only for principle-level disputes.

## Ownership shift — who owns a classification change

PRE-BUILDER

CPA opens ticket → triage → engineer writes rule → PR → review → ship.

3–5 days

PER CHANGE

POST-BUILDER

CPA opens Builder → preview against last 30d → save. Engineer not in loop.

~15 min

PER CHANGE

↓85%

ENGINEERING DEPENDENCY FOR CLASSIFICATION-LOGIC CHANGES

↓ 85% isn't "design replaced engineering." It's "design redirected engineering to where they actually moved the needle."

## Three artifacts that proved the shift wasn't cosmetic

PROOF · NOT JUST A PROCESS DIAGRAM

ARTIFACT · 01

### Principles doc as tiebreaker.

The day the team referenced the principles doc in a PR comment thread to settle a design/eng debate. I was out. They didn't need me.

ARTIFACT · 02

### First CPA-authored rule shipped.

First rule that went live through the Builder without engineering review. The Builder was the safety net — preview + conflict detection — not the code review.

ARTIFACT · 03

### "Should this still come to us?"

Rule Builder PR thread where engineering explicitly asked whether classification logic changes should still route through them. Answer: no. The handoff moment.

# The four numbers I'll defend.

Each paired with the principle it most directly proves. Each sourced and caveated.

↓40%

ACCOUNTING PROCESSING TIME

v1 baseline 6h avg/client · v2 3.6h avg/client

P1 · TRUST-FIRST · PILOT 4 FIRMS 6 WKS

↓45%

CPA ONBOARDING TIME

v1 median 14.5d · v2 median 8.0d · n=22 vs 22

P2 · TRANSPARENCY · FIRM-SIZE CONTROLLED

↑32%

DEMO → CUSTOMER CONVERSION

v1 baseline 28% · v2 37% · n=64 demos

P3 · AUGMENTATION · CONFOUNDED – SEE CAVEAT

↓85%

ENGINEERING DEPENDENCY

Rule-change cycles: eng → CPA-authored

SYSTEMS OUTCOME · RULE-BUILDER SHIPPED WK 26

| METRIC                        | SOURCE                                                       | CAVEAT                                                                                              | STRENGTH |
|-------------------------------|--------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|----------|
| • ↓40% processing time        | Pilot timesheet capture · 4 firms · 6 wk pre/post            | Self-reported, cross-checked with 6 shadow sessions. Pilot firms higher-volume than median.         | STRONG   |
| • ↓45% onboarding time        | v1 cohort n=22 vs v2 cohort n=22 · firm-size band controlled | Controlled for firm-size band, not same firms. Methodology on Slide 23.                             | MODERATE |
| • ↑32% demo conversion        | Sales pipeline · v1 28% (n=32) vs v2 37% (n=32)              | <b>Confounded with concurrent sales-team training change.</b> Directional. Attribution on Slide 21. | STRONG   |
| • ↓85% engineering dependency | Rule-change PRs · 12wk pre-Builder vs 12wk post              | PR count, not eng-hours. Hours saved correlates but isn't 1:1.                                      | MODERATE |

## Two honest misses.

### MISS · 01

#### Shipped Rule Builder without preview mode for 4 weeks.

**Why I missed:** Optimized for speed of authoring. Assumed CPAs would test rules by triggering them on real data and rolling back.

**What changed:** Two CPAs broke their classifications in week 18. Preview-mode shipped week 22 ("Live test, last 30 days" — Slide 25).

***The miss generated the principle.*** "Preview before commit" got added to the deck as a P2 sub-principle.

### MISS · 02

#### Confidence-tier thresholds were guessed for 7 weeks.

**Why I missed:** Picked 70/40 as "feels right" thresholds in week 5. Didn't bin against override-rate empirically.

**What changed:** Week 12 analysis revealed elbows at 60% and 90%. Re-tiered.

*Echo of Slide 6 quote — "How do I defend this to the IRS if I can't explain the logic?" — landed twice in pilot. Worth the deliberate callback.*

BACKUP PACK

# Seven slides anticipating likely hiring-manager objections.

Pulled in only if a question opens the thread. Each is sourced.

B

# Personas — David Chen and Mya.

Derived from 12 interviews + 6 shadow sessions. Composites; not real individuals.

PERSONA · 01

☰

David Chen — Senior CPA

PRIMARY

- **Role:** Partner; signs returns; license on the line.
- **Goals:** Defend every number to the IRS. Move quickly through known patterns.
- **Distrusts:** Surprises — "why is this different from last year?"
- **Meets product at:** Audit drawer, exports, anomaly review.
- **Doesn't use:** Rule Builder (delegates to associates).

**Touches:** audit drawer, exports, anomalies, review queue.

**Designing for David** = audit drawer + exports = P1 (trust) + P2 (transparency) territory.

PERSONA · 02

⋮

Mya — Tax Analyst

SECONDARY

- **Role:** Associate; high-volume; reviews 20–40 transactions per session.
- **Goals:** Get through the queue. Flag anything David needs to see.
- **Distrusts:** Anything that hides her keystroke count.
- **Meets product at:** Transactions table, override flow, Rule Builder.
- **Power-uses:** Bulk-select (gated), keyboard shortcuts, saved filters.

**Touches:** transactions, override, Rule Builder, bulk.

**Designing for Mya** = transactions + override = could be any principle, lives in P3 (augmentation) most often.

# Attribution — separating design from confounds.

For the ↑32% demo conversion specifically — the metric with the heaviest caveat.

The **qualitative thesis**: the 78% is the strongest stand-alone signal that design — not training, not category tailwind — moved the conversion number. But naming the confounds is the case *for* transparency, applied to my own metrics.

78%

PROCUREMENT CALLS

of post-v2 procurement calls cited **transparency / audit defensibility** as a primary reason for adoption (n=18).

Naming the confounds — and the limits of what can be cleanly attributed — *is* the case for transparency, applied to my own metrics.

## Attribution breakdown

| FACTOR                                | ~%  | TYPE     |
|---------------------------------------|-----|----------|
| • Transparency / audit drawer         | 40% | DESIGN   |
| • Rule Builder                        | 20% | DESIGN   |
| Sales-team training change            | 15% | CONFOUND |
| Category tailwind (crypto-tax demand) | 15% | CONFOUND |
| Other / unattributable                | 10% | MIXED    |

SOURCE: hand-coded from 18 post-call notes · sales-team conversion rationale tagging · pilot Q1–Q2

# Counterfactual — what 3 engineers instead of 2 would have bought.

## What was descope

| ITEM                    | WHERE CUT | COST                  |
|-------------------------|-----------|-----------------------|
| Bulk override           | Week 14   | Shipped 10 wks late   |
| Multi-firm rule sharing | Week 16   | Deferred to v3        |
| Mobile review           | Week 8    | Deferred to v3        |
| Real-time collab        | Week 10   | Deferred indefinitely |
| Advanced anomaly ML     | Week 18   | Shipped rule-based v1 |

## What I'd protect with more capacity

**Wouldn't accelerate:** the principles, the research cadence, the v1 → v2 pivot timing. More capacity wouldn't have surfaced findings faster — pilot-firm scheduling was the bottleneck.

**Would accelerate:** Rule Builder preview-mode (shipped wk 22 → could've shipped wk 18). Multi-firm rule sharing as the v2 leverage compounder.

FOUNDING-DESIGNER LESSON

*More capacity buys system depth, not better principles. Knowing which is which is part of the role.*

# How the onboarding metric was measured.

## Method

- **Cohort A (v1):** 22 CPAs onboarded between Sep 2024 – Dec 2024 on v1 product.
- **Cohort B (v2):** 22 CPAs onboarded between Feb 2025 – Apr 2025 on v2 product.
- **Matched on:** firm-size band ( $\leq 10$  / 11–50 / 51+), prior crypto-tax tool experience (yes/no).
- **Metric:** median days from sign-up to first client-load completion.

14.5d → 8.0d

↓45% MEDIAN ONBOARDING

v1 median 14.5d · v2 median 8.0d · n=22 vs n=22

## Caveats — what this doesn't capture

- **Different time windows.** v1 cohort onboarded pre-tax-season; v2 cohort during tax season — different urgency baseline.
- **Same firms? No.** Different individuals; controlled for firm-size band, not for individual experience level.
- **Onboarding includes sales onboarding.** A sales-process change may have contributed.
- **Median, not mean.** Tail was long in both cohorts (one v1 user took 45d).

A 45% number with named caveats holds up better than one without — internally and in front of customers.


# Competitive matrix — Computis vs. the field.

Marketing-page snapshot, Mar 2025. Not exhaustive — coded from public docs + 300+ G2 reviews.

**Strategic position — what the matrix earns:** The two unique features — no-code rules and anomaly detection — both follow from the same insight that broke v1: *the rejection wasn't AI, it was opacity*. Competitors automate; Computis exposes the controls.

| FEATURE                 | TAXBIT | CRYPTIO | COINTRACKER | COMPUTIS  |
|-------------------------|--------|---------|-------------|-----------|
| Auto-classify           | ✓      | ✓       | ✓           | ✓         |
| Override control        | ✗      | ✗       | 🕒           | REQUIRED  |
| Audit trail (IRS-ready) | 🕒      | ✗       | ✗           | IRS-READY |
| No-code rule authoring  | ✗      | ✗       | ✗           | UNIQUE    |
| Anomaly detection       | ✗      | 🕒       | ✗           | UNIQUE    |

COMPETITOR COMPLAINT ANALYSIS • CODED FROM 300+ G2 REVIEWS




WHAT THIS TELLS US

**82% of competitor complaints** are opacity + control issues — not accuracy or speed. The two unique columns in the matrix aren't features. They're the answers to that 82%.

Sources: G2 Reviews, Capterra, TrustRadius • Q4 2024 • n=300+ reviews coded for theme • competitive feature mapping from public marketing pages, Mar 2025

# Project timeline — 10 months, ~60% of decisions in 6 weeks.

Decision density was front-loaded around the v1→v2 pivot. Same 40-week rail shown on Slide 8 — annotated here for what shipped where.



## Decision density

Roughly **60% of decisions that shipped** were made in the six weeks between the joint session (wk 7) and the v2 design freeze (wk 12). The first six weeks earned the right.

## Why the cadence held

Two-week design / one-week eng-feedback. The v1→v2 pivot didn't require a re-plan because the cadence was already tight enough to absorb a redirect inside one sprint.

FOUNDING-DESIGNER LESSON

Compressed timelines reward research-as-leverage. Six weeks of evidence collected before any UI decisions meant the second six weeks could be argued in hours, not days. *Speed is downstream of clarity.*

# Three threads I'd pull next.

Roadmap fragments. Each ties back to a specific finding or miss in the deck.

THREAD · 01

## Confidence calibration ↔ override-rate feedback loop.

Today's tiers are static thresholds; should evolve with overrides per firm. Closes the loop between Slide 11 (tier system) and Slide 14 (override taxonomy).

THREAD · 02

## Multi-CPA review surfaces.

Pilot was single-CPA; larger firms have second-review workflows we didn't model. David delegates to Mya — the surface doesn't know that.

THREAD · 03

## Rule Builder as marketplace.

Firm-authored rules currently silo; CPA-to-CPA sharing is the next leverage point. The v2 compounder Slide 22 named.

## Where this design system goes next.

Computis becomes the IRS-defensible AI standard — the platform other tax-tooling vendors benchmark their explainability against.

MOVE • 01

### Multi-firm contributor model.

The Rule Builder's per-rule preview already supports it. What's missing: contributor permissions layer + rule-publishing flow.

EST. COST • ~6 wk design • ~8 wk eng

MOVE • 02

### Cross-product reuse.

Confidence-tier UI + audit drawer are the most portable artifacts. Same components, different shell — anchor a CPA-firm dashboard inside QuickBooks or a partner surface.

ENABLES • partner-product surface area

MOVE • 03

### AI-explanation pattern as platform feature.

Audit drawer is a Computis surface today. Year 2: an SDK component vendors pay for. The explainability layer the tax industry licenses.

SHIFTS • product → platform

**What this earns competitively:** from "Computis is an AI tax tool" to "Computis is the AI-explainability layer the tax industry depends on." Same code, larger surface area, harder to copy.

# Three principles. One system. Two misses I'd answer for. Four numbers I'll defend.

Computis was a study in **earning AI trust from a license-on-the-line audience**.  
Augmentation over automation. Transparency under regulation. A design  
system that survived a ten-month run with one designer and two engineers.

P1 • TRUST-FIRST

P2 • TRANSPARENCY

P3 • AUGMENTATION

SEAN SMITH • FOUNDING PRODUCT DESIGNER

smithdesign.live • computis.netlify.app

## THE DECK IN ONE FRAME

↓40%

ACCOUNTING PROCESSING TIME

↓45%

CPA ONBOARDING TIME

↑32%

DEMO → CUSTOMER CONVERSION

↓85%

ENG DEPENDENCY

## TWO MISSES I'D ANSWER FOR

**Rule Builder shipped without preview** — 4 weeks late on the fix.

**Tier thresholds were guessed for 7 weeks** — re-tiered on data at wk 12.